

iris =

	sepal_length	sepal_width	petal_length	petal_width	species
1	5.1	3.5	1.4	0.2	"Iris-setosa"
2	4.9	3.0	1.4	0.2	"Iris-setosa"
3	4.7	3.2	1.3	0.2	"Iris-setosa"
4	4.6	3.1	1.5	0.2	"Iris-setosa"
5	5.0	3.6	1.4	0.2	"Iris-setosa"
6	5.4	3.9	1.7	0.4	"Iris-setosa"
7	4.6	3.4	1.4	0.3	"Iris-setosa"
8	5.0	3.4	1.5	0.2	"Iris-setosa"
9	4.4	2.9	1.4	0.2	"Iris-setosa"
10	4.9	3.1	1.5	0.1	"Iris-setosa"
more					
150	5.9	3.0	5.1	1.8	"Iris-virginica"

```
iris = @pipe CSV.read(download("https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"), DataFrame, header=["sepal_length", "sepal_width", "petal_length", "petal_width", "species"]) |> coerce!(_, :species => Multiclass)
```

names	scitypes	types
sepal_length	Continuous	Float64
sepal_width	Continuous	Float64
petal_length	Continuous	Float64
petal_width	Continuous	Float64
species	Multiclass{3}	CategoricalValue{String15, UInt32}

```
schema(iris)
```

0

```
isnothing.(iris) |> Matrix |> sum
```

```
y_col = "species"
```

```
y_col = "species"
```

splitting the dataset

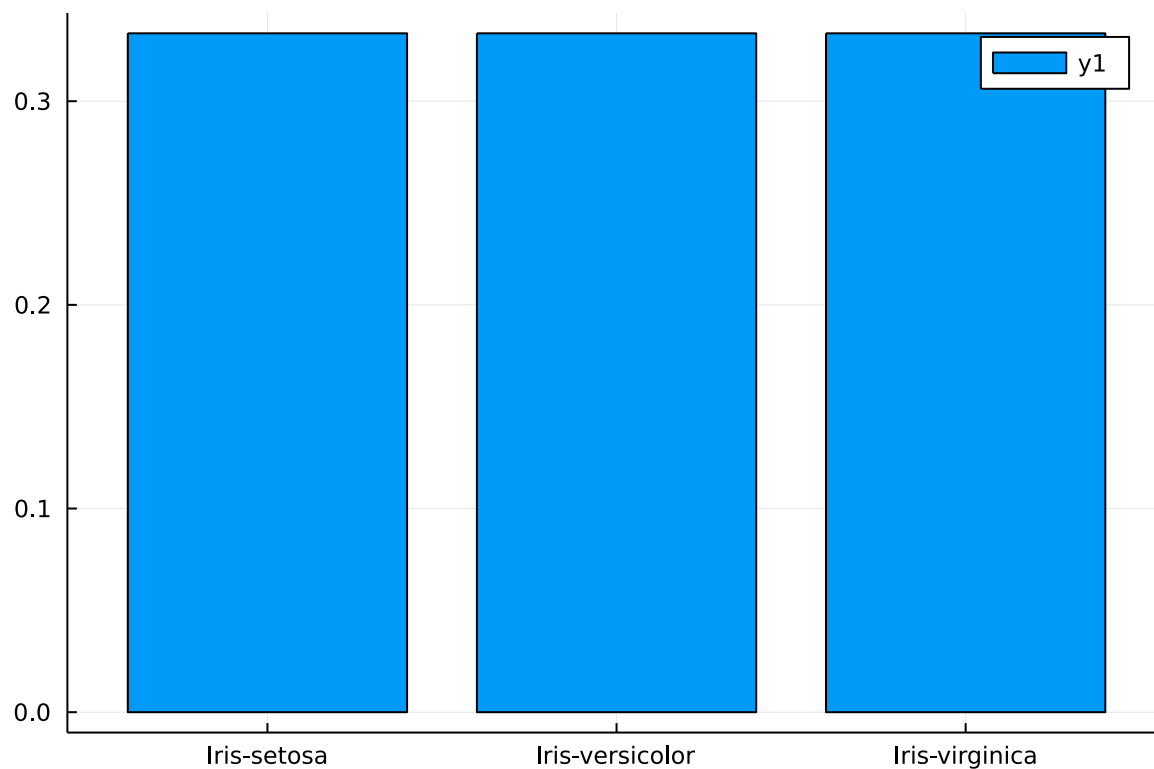
```
train, test = stratifiedobs(row->row[:species], iris); nothing
```

we don't look at testset

	variable	sepal_length	sepal_width	petal_length	petal_width	specie
1	"mean"	5.85143	3.0181	3.76095	1.19143	nothing
2	"std"	0.859686	0.442411	1.795	0.754488	nothing
3	"min"	4.3	2.0	1.0	0.1	"Iris-setosa"
4	"q25"	5.1	2.7	1.5	0.3	nothing
5	"median"	5.8	3.0	4.4	1.3	nothing
6	"q75"	6.4	3.3	5.1	1.8	nothing
7	"max"	7.9	4.2	6.9	2.5	"Iris-virginica"
8	"nunique"	nothing	nothing	nothing	nothing	3
9	"nmissing"	0	0	0	0	0
10	"first"	5.5	4.2	1.4	0.2	"Iris-setosa"
11	"last"	5.4	3.7	1.5	0.2	"Iris-setosa"
12	"eltype"	Float64	Float64	Float64	Float64	CategoricalValue{St

```
• @pipe describe(train, :all) |> permutedims(_, 1)
```

but, it's okay to look at the population distribution of the entire dataset(including testset)

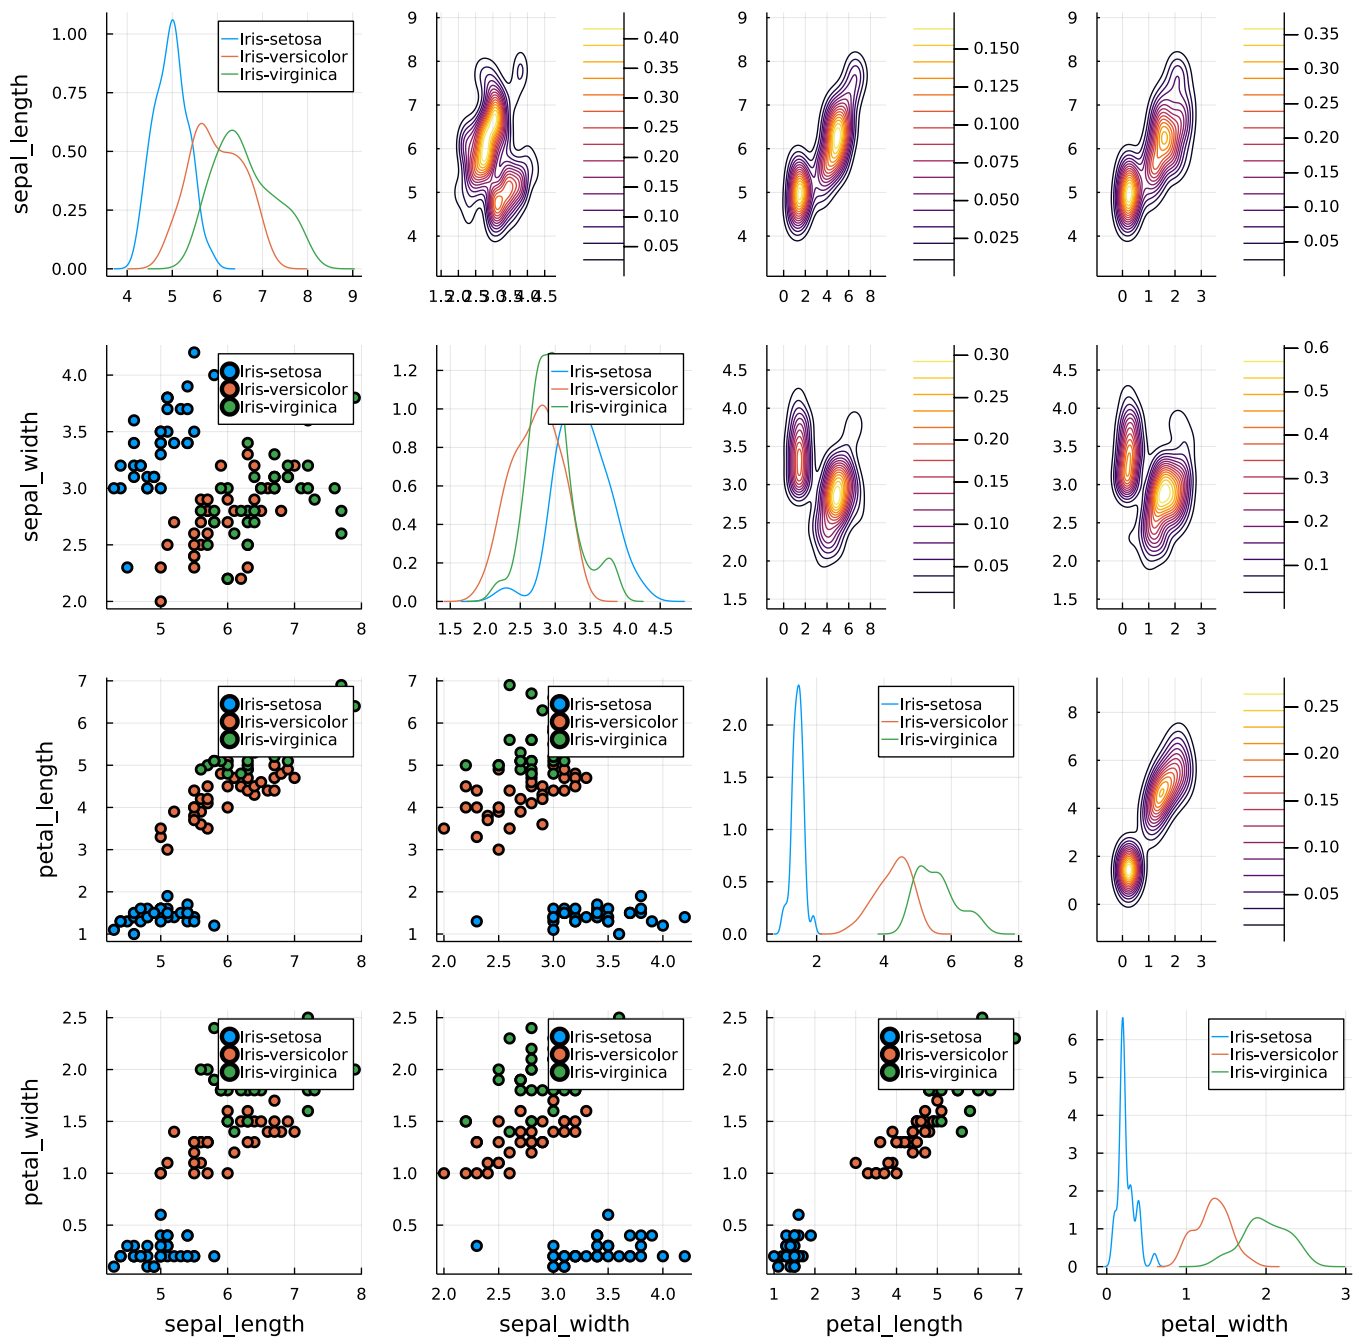


```
• let
•   species_dist = countmap(iris.species)
•   species_dist = convert(Dict{eltype(keys(species_dist)), Float64}, species_dist)
•   for (species, count) ∈ species_dist
•       species_dist[species] /= size(iris, 1)
•   end
•   bar(species_dist)
• end
```

population is normally distributed: no target imbalance

# EDA

---



```

• begin
•   features = names(train)[1:end-1]
•   n = length(features)
•   ps = fill(Plots.plot(), n, n)
•
•   for i ∈ 1:n, j ∈ 1:n
•     x, y = train[:, features[i]], train[:, features[j]]
•     species = train[:, :species]
•     ps[i, j] = if j < i
•       dens = KernelDensity.kde((x, y))
•       Plots.plot(dens)
•     elseif j > i
•       Plots.plot(x, y, seriestype=:scatter, group=species)
•     else
•       Plots.plot(x, y, seriestype=:density, group=species)
•     end
•   end

```

```
•         i == 1 && Plots.ylabel!(ps[i, j], features[j])  
•         j == n && Plots.xlabel!(ps[i, j], features[i])  
•     end  
•     Plots.plot(ps..., layout = (n, n), size=(1000,1000))  
• end
```

Assignment to `y` in soft scope is ambiguous because a global variable by the same name exists: `y` will be treated as a new local. Disambiguate by using `local y` to suppress this warning or `global y` to assign to the existing global variable.

we find that:


- can't see any significant outliers that would skew the distributions
- Because Petal length and Petal width have separable distributions for different iris species, these 2 variables would have higher influence over classifying the species.
- We can also observe from the scatterplots of *petal length* x *petal width*, *petal width* x *sepal width* and *petal width* x *sepal length*, that the species clusters are clearly distinguishable and can easily be separated by a line.
- Hence, a single linear model(s) would be best suited for this type of classification here, without requiring of any feature transformations as the features are already distinguishable.
- The setosa species is the most easily distinguishable because of its small feature size.

plot\_heatmap (generic function with 1 method)



```
@pipe cor(Matrix(train[:, Not(y_col)])) |> plot_heatmap(_, xlabel=names(train)[1:end-1])
```

- if we look further we can see that petal\_width is highly +vely correlated to petal\_length and sepal\_length. Similarly, petal\_length and sepal\_length are also significantly correlated.
- Also, sepal\_width is slightly -vely correlated with other variables.



```
• @pipe cov(Matrix(train[:, Not(y_col)])) |> plot_heatmap(_, xlabel=names(train)[1:end-1])
```

A strong correlation is present between petal width and petal length.

## Modeling

we earlier found out that the dataset is not too complex and simple models(most probably linear models) would be better suited for this problem.

```
((105, 4), (105), (45, 4), (45))
```

```
• begin
•   X, y = iris[:, Not(y_col)], iris[:, y_col]
•   (X_train, y_train), (X_test, y_test) = stratifiedobs((X, y))
•   size(X_train), size(y_train), size(X_test), size(y_test)
• end
```

```
LogisticClassifier = MLJLinearModels.LogisticClassifier
```

```
• LogisticClassifier = @load LogisticClassifier pkg=MLJLinearModels add=true
```

```
import MLJLinearModels ② verbosity=0`.
```

```
pipe = Pipeline431(
  logistic_classifier = LogisticClassifier(
    lambda = 1.0,
    gamma = 0.0,
    penalty = :l2,
    fit_intercept = true,
    penalize_intercept = false,
    scale_penalty_with_samples = true,
    solver = nothing))
```

- `pipe = @pipeline LogisticClassifier`

```
model =
```

```
Machine trained 0 times; caches data
```

```
model: Pipeline431(logistic_classifier = LogisticClassifier(lambda = 1.0, ...))
```

```
args:
```

```
1: Source @863 ↗ `ScientificTypesBase.Table{AbstractVector{ScientificTypesBase.Contin
2: Source @847 ↗ `AbstractVector{ScientificTypesBase.Multiclass{3}}`
```

- `model = machine(pipe, X_train, y_train)`

```
PerformanceEvaluation object with these fields:
```

```
measure, operation, measurement, per_fold,
per_observation, fitted_params_per_fold,
report_per_fold, train_test_rows
```

```
Extract:
```

measure	operation	measurement	1.96*SE	per_fold	...
Accuracy()	predict_mode	0.827	0.0577	[0.833, 0.83 ...	
LogLoss( tol = 2.220446049250313e-16)	predict	0.693	0.0256	[0.682, 0.71 ...	...

1 column omitted

- `MLJ.evaluate!(model, resampling=CV(nfolds=6, rng=StableRNG(32)),`
- `measures=[MLJ.accuracy, log_loss])`

```
Evaluating over 6 folds: 33%[=====>] ETA: 0:00:0
4 [KEvaluating over 6 folds: 50%[=====>] ETA: 0:00:0
3 [KEvaluating over 6 folds: 100%[=====] Time: 0:00:0
7 [K
```

## Model Tuning

```
ranges =
```

```
[NumericRange(0.0 ≤ logistic_classifier.lambda ≤ 1.0; origin=0.5, unit=0.5) on log scale,
```

- `ranges = [range(pipe, :(logistic_classifier.lambda), lower=0.0, upper=1.0,`
- `scale=:log),`
- `range(pipe, :(logistic_classifier.gamma), lower=0.0, upper=1.0, scale=:log),`
- `range(pipe, :(logistic_classifier.penalty), values=[:l2, :l1, :en, :none]),`
- `]`





```

Pipeline431(
  logistic_classifier = LogisticClassifier(
    lambda = 0.37919381454429746,
    gamma = 0.6722497885235214,
    penalty = :none,
    fit_intercept = true,
    penalize_intercept = false,
    scale_penalty_with_samples = true,
    solver = nothing))

```

```
• fitted_params(tuned_model).best_model
```

```
r =
```

```

(best_model = Pipeline431(
  logistic_classifier = LogisticClassifier(
    lambda = 0.379194,
    gamma = 0.67225,
    penalty = :none,
    fit_intercept = true,
    penalize_intercept = false,
    scale_penalty_with_samples = true,
    solver = nothing))
, best_history_entry = (model =

```

```
• r = report(tuned_model)
```

```
0.9230769230769231
```

```
• r.best_history_entry.measurement[1]
```

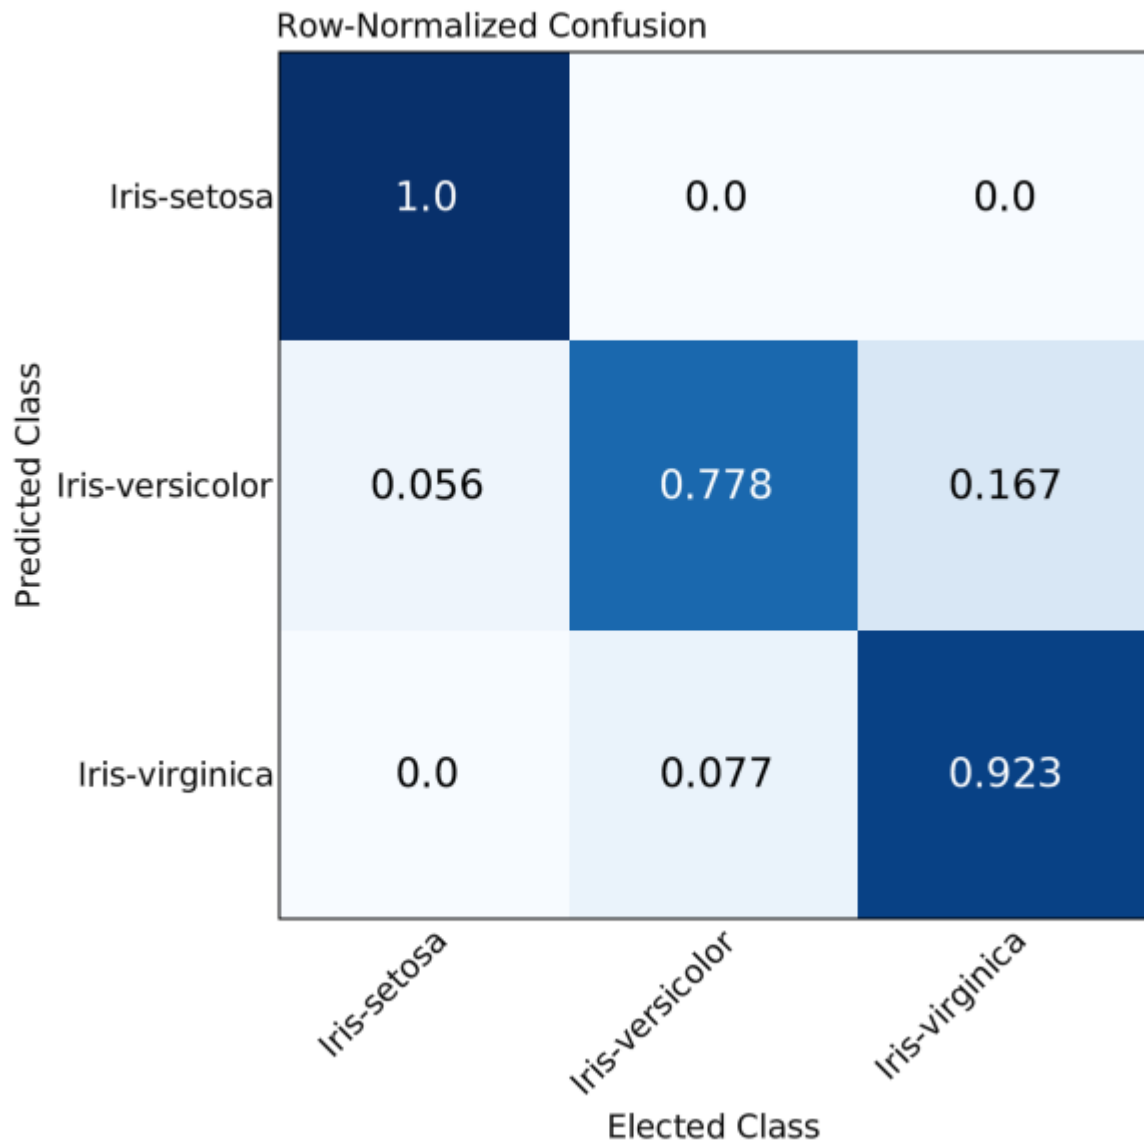
```
 $\hat{y}$  =
```

```
CategoricalArrays.CategoricalVector{InlineStrings.String15, UInt32, InlineStrings.String1
```

```
•  $\hat{y}$  = MLJ.predict(tuned_model, X_test) .|> mode
```

```
0.8888888888888888
```

```
• Accuracy()( $\hat{y}$ , y_test)
```



```
• begin
•   cnf_mat = ConfusionMatrix()(ŷ, y_test)
•   Lighthouse.plot_confusion_matrix(cnf_mat.mat, cnf_mat.labels, :Row)
• end
```

The classes are un-ordered,  
using order: InlineStrings.String15["Iris-setosa", "Iris-versicolor", "Iris-virginica"].  
To suppress this warning, consider coercing to OrderedFactor.

todo: classification report

```
• WGLMakie.activate!()
```

```
• begin
•   using Pkg
•   Pkg.activate(pwd())
•   # Pkg.add(["CSV", "DataFrames"])
•   # Pkg.add("StatsBase")
•   # Pkg.add("Pipe")
•   # Pkg.add("MLDataUtils")
•   # Pkg.add("MLJ")
•   # Pkg.add("StatsPlots")
•   # Pkg.add("KernelDensity")
•   # Pkg.add("WGLMakie")
•   # # Pkg.rm("WGLMakie")
•   # Pkg.add("StableRNGs")
•   # Pkg.add(name="Lighthouse", version="0.14")
•   # Pkg.add("PlutoUI")
•   # Pkg.update()
•
•   using Plots, StatsPlots, PlutoUI
•   using KernelDensity
•   using WGLMakie
•   using DataFrames, CSV
•   using Statistics, StatsBase, StableRNGs
•   using Pipe
•   using MLDataUtils, MLJ, Lighthouse
• end
```

Activating project at `~/Documents/my-portfolio-projects/machine learning projects/iris-flowers`

